

## Kubernetes\_3 notes

This one manifest file will create both service and pod  
We already demonstrated LoadBalancer service

### Kubernetes NodePort service

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: NodePort
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~$ vi k8s-pod-svc-manifest-NodePort.yml
ubuntu@ip-172-31-9-165:~$ cat k8s-pod-svc-manifest-NodePort.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: NodePort
```

```

selector:
  app: javawebapp
ports:
  - port: 80
    targetPort: 8080
...

```

```


ubuntu@ip-172-31-9-165:~$ ls -l
total 34160
drwxr-xr-x 3 ubuntu ubuntu 4096 May 16 18:46 aws
-rw-rw-r-- 1 ubuntu ubuntu 34958926 May 17 23:42 eksctl.tar.gz
-rw-rw-r-- 1 ubuntu ubuntu 458 May 18 22:39 k8s-pod-manifest-new.yml
-rw-rw-r-- 1 ubuntu ubuntu 229 May 18 21:29 k8s-pod-manifest.yml
-rw-rw-r-- 1 ubuntu ubuntu 457 May 24 14:50 k8s-pod-svc-manifest-NodePort.yml
-rw-rw-r-- 1 ubuntu ubuntu 195 May 18 22:05 k8s-service-manifest.yml
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml

```

```

eksctl create cluster --name my-eks-cluster --region ca-central-1 --node-type t2.medium --zones ca-central-1a,ca-central-1b

```

|                          |                                 |                     |         |                                                                                                                                                                             |    |
|--------------------------|---------------------------------|---------------------|---------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| <input type="checkbox"/> | my-eks-cluster-ng-ee43d2d8-Node | i-0f5dfab3fdbb83127 | Running |       | t2 |
| <input type="checkbox"/> | EKS-host                        | i-01289fc5ca918b25f | Running |     | t2 |
| <input type="checkbox"/> | my-eks-cluster-ng-ee43d2d8-Node | i-0a3cb31674ad30197 | Running |   | t2 |

```

ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created

```

```

ubuntu@ip-172-31-9-165:~$ kubectl get pod
NAME      READY  STATUS   RESTARTS  AGE
javawebapp 1/1    Running  0         75s

```

```

ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get pod
NAME      READY  STATUS   RESTARTS  AGE
javawebapp 1/1    Running  0         75s
ubuntu@ip-172-31-9-165:~$ kubectl get service
NAME      TYPE      CLUSTER-IP  EXTERNAL-IP  PORT(S)    AGE
javawebappsvc NodePort  10.100.17.144 <none>      80:30138/TCP 108s
kubernetes ClusterIP 10.100.0.1   <none>      443/TCP     9m42s

```

```
2025-05-24 15:24:01 [✓] EKS cluster "my-eks-cluster" in "ca-central-1" region i
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get pod
NAME          READY   STATUS    RESTARTS   AGE
javawebapp    1/1     Running   0           75s
ubuntu@ip-172-31-9-165:~$ kubectl get service
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
javawebappsvc NodePort     10.100.17.144 <none>        80:30138/TCP     108s
kubernetes    ClusterIP    10.100.0.1     <none>        443/TCP          9m42s
ubuntu@ip-172-31-9-165:~$
```

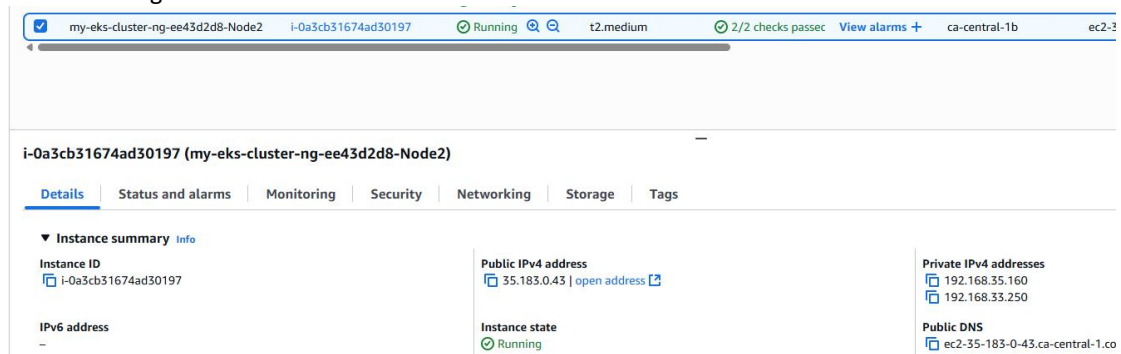
```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                                     NOMINATED
NODE READINESS GATES
javawebapp    1/1     Running   0           36m   192.168.40.10 ip-192-168-35-160.ca-central-
1.compute.internal <none>    <none>
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get service
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
javawebappsvc NodePort     10.100.17.144 <none>        80:30138/TCP     108s
kubernetes    ClusterIP    10.100.0.1     <none>        443/TCP          9m42s
```

A NodePort# will be assigned automatically

Using NodePort we can access our app with WorkerNode public IP address

Pod is running in Node2



my-eks-cluster-ng-ee43d2d8-Node2 i-0a3cb31674ad30197 Running t2.medium 2/2 checks passed View alarms + ca-central-1b ec2-3

i-0a3cb31674ad30197 (my-eks-cluster-ng-ee43d2d8-Node2)

Details Status and alarms Monitoring Security Networking Storage Tags

▼ Instance summary info

Instance ID  
i-0a3cb31674ad30197

Public IPv4 address  
35.183.0.43 | open address

Private IPv4 addresses  
192.168.35.160  
192.168.33.250

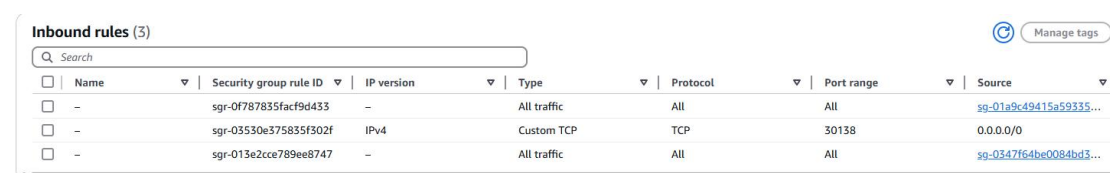
IPv6 address  
-

Instance state  
Running

Public DNS  
ec2-35-183-0-43.ca-central-1.co

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                                     NOMINATED
NODE READINESS GATES
javawebapp    1/1     Running   0           36m   192.168.40.10 ip-192-168-35-160.ca-central-
1.compute.internal <none>    <none>
```

So we Add 30138 to Inbound Rules of Node2

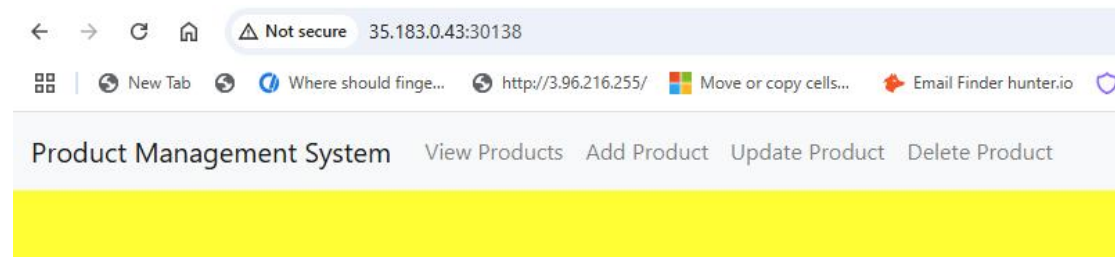


Inbound rules (3)

Search

| Name | Security group rule ID | IP version | Type        | Protocol | Port range | Source                 |
|------|------------------------|------------|-------------|----------|------------|------------------------|
| -    | sgr-0f787835facf9d433  | -          | All traffic | All      | All        | sg-01a9c49415a59335... |
| -    | sgr-03530e375835f302f  | IPv4       | Custom TCP  | TCP      | 30138      | 0.0.0.0/0              |
| -    | sgr-013e2cce789ee8747  | -          | All traffic | All      | All        | sg-0347f64be0084bd3... |

We pick the Public IP then add Port#  
<http://35.183.0.43:30138/>



NodePort can also be used to expose Application Pods

```
ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "javawebapp" deleted
service "javawebappsvc" deleted
service "kubernetes" deleted
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
javawebappsvc NodePort      10.100.206.94 <none>       80:30447/TCP     48s
kubernetes    ClusterIP    10.100.0.1    <none>       443/TCP          119s
```

This time it is a different port number

```
javawebappsvc NodePort 10.100.206.94 <none> 80:30447/TCP 48s
```

```
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "javawebapp" deleted
service "javawebappsvc" deleted
service "kubernetes" deleted
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
kubernetes    ClusterIP    10.100.0.1    <none>       443/TCP          28s
ubuntu@ip-172-31-9-165:~$ ls -l
total 34160
drwxr-xr-x 3 ubuntu ubuntu    4096 May 16 18:46 aws
-rw-rw-r-- 1 ubuntu ubuntu 34958926 May 17 23:42 eksctl.tar.gz
-rw-rw-r-- 1 ubuntu ubuntu    458 May 18 22:39 k8s-pod-manifest-new.yml
-rw-rw-r-- 1 ubuntu ubuntu    229 May 18 21:29 k8s-pod-manifest.yml
-rw-rw-r-- 1 ubuntu ubuntu    454 May 24 15:09 k8s-pod-svc-manifest-NodePort.yml
-rw-rw-r-- 1 ubuntu ubuntu    195 May 18 22:05 k8s-service-manifest.yml
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
javawebappsvc NodePort      10.100.206.94 <none>       80:30447/TCP     48s
kubernetes    ClusterIP    10.100.0.1    <none>       443/TCP          119s
ubuntu@ip-172-31-9-165:~$
```

NodePort#:

If we don't specify NodePort number in service-manifest yaml file, then Kubernetes will assign a random NodePort number for the service within the range 30000 to 32767. However, we can specify NodePort number in service-manifest yaml

kubectl apply -f <yaml> ---> Access the application using Public IP of workerNode alongside NodePort number. Also add NodePort number in Security group

Add one extra line

nodePort: 30044

---

apiVersion: v1

kind: Pod

metadata:

name: javawebapp

labels:

app: javawebapp

spec:

containers:

- name: javawebappcontainer

image: hacker123shiva/springbt-in-docker:latest

ports:

- containerPort: 8080

---

apiVersion: v1

kind: Service

metadata:

name: javawebappsvc

spec:

type: NodePort

selector:

app: javawebapp

ports:

- port: 80

targetPort: 8080

nodePort: 30044

...

ubuntu@ip-172-31-9-165:~\$ vi k8s-pod-svc-manifest-NodePort.yml

ubuntu@ip-172-31-9-165:~\$ cat k8s-pod-svc-manifest-NodePort.yml

---

apiVersion: v1

kind: Pod

metadata:

name: javawebapp

labels:

app: javawebapp

spec:

containers:

- name: javawebappcontainer

image: hacker123shiva/springbt-in-docker:latest

ports:

- containerPort: 8080

---

apiVersion: v1

kind: Service

metadata:

```

name: javawebappsvc
spec:
  type: NodePort
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30044
...

```

```

...
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-NodePort.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
javawebapp    1/1     Running   0           8s
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME            TYPE          CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
javawebappsvc   NodePort      10.100.9.42    <none>        80:30044/TCP     12s
kubernetes      ClusterIP     10.100.0.1     <none>        443/TCP          99s
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                                     NOMINATED
NODE READINESS GATES
javawebapp    1/1     Running   0           77s   192.168.43.212 ip-192-168-35-160.ca-central-1.compute-1.internal <none> <none>

```

192-168-35-160 is in Node2

my-eks-cluster-ng-ee43d2d8-Node2 i-0a3cb31674ad30197 Running t2.medium 2/2 checks passed View alarms + ca-central-1b ec2-35-183

**i-0a3cb31674ad30197 (my-eks-cluster-ng-ee43d2d8-Node2)**

**Instance summary** info

Instance ID: i-0a3cb31674ad30197

IPv6 address: -

Hostname: fupa

Public IPv4 address: 35.183.0.43 | open address

Instance state: Running

Private IP DNS name (IPv4 only):

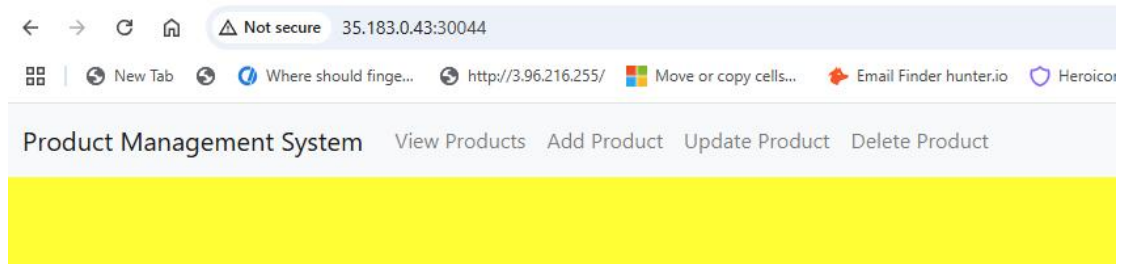
Private IPv4 addresses: 192.168.35.160, 192.168.33.250

Public DNS: ec2-35-183-0-43.ca-central-1.compute

## Security groups

| Inbound rules (4)        |      |                        |            |             |          |            |  |
|--------------------------|------|------------------------|------------|-------------|----------|------------|--|
| Search                   |      |                        |            |             |          |            |  |
| <input type="checkbox"/> | Name | Security group rule ID | IP version | Type        | Protocol | Port range |  |
| <input type="checkbox"/> | -    | sgr-0f787835facf9d433  | -          | All traffic | All      | All        |  |
| <input type="checkbox"/> | -    | sgr-0b93f63a159b4af4a  | IPv4       | Custom TCP  | TCP      | 30044      |  |
| <input type="checkbox"/> | -    | sgr-03530e375835f302f  | IPv4       | Custom TCP  | TCP      | 30138      |  |
| <input type="checkbox"/> | -    | sgr-013e2cce789ee8747  | -          | All traffic | All      | All        |  |

<http://35.183.0.43:30044/>



```
ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "javawebapp" deleted
service "javawebappsvc" deleted
service "kubernetes" deleted
```

Cluster IP service

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
---
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: ClusterIP
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
---
```

```
ubuntu@ip-172-31-9-165:~$ vi k8s-pod-svc-manifest-clusterIP.yml
ubuntu@ip-172-31-9-165:~$ cat k8s-pod-svc-manifest-clusterIP.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
```

```
- containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: ClusterIP
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-clusterIP.yml
pod/javawebapp created
service/javawebappsvc created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-clusterIP.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
javawebapp 1/1     Running   0           30s
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME            TYPE       CLUSTER-IP   EXTERNAL-IP  PORT(S)   AGE
javawebappsvc   ClusterIP  10.100.143.74 <none>       80/TCP    35s
kubernetes      ClusterIP  10.100.0.1    <none>       443/TCP   33m
```



```

ubuntu@ip-172-31-9-165:~$ cat k8s-pod-svc-manifest-clusterIP.yml
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: ClusterIP
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-manifest-clusterIP.yml
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
javawebapp    1/1     Running   0           30s
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
javawebappsvc ClusterIP    10.100.143.74 <none>        80/TCP      35s
kubernetes    ClusterIP   10.100.0.1    <none>        443/TCP     33m
ubuntu@ip-172-31-9-165:~$

```

ClusterIP will be used to access our pods within Cluster and one Static IP will be created to access the pods

ClusterIP cannot be access from outside network. So NodePort is required

```

ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "javawebapp" deleted
service "javawebappsvc" deleted
service "kubernetes" deleted

```

Kubenetes Namespaces:

Why we need Namespace? To group the resources

frontend-app-pods --> frontend-app-ns

backend-app-pods --> backend-app-ns

database-pods --> database-ns

```

ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME          STATUS   AGE
default       Active   4h19m
kube-node-lease Active   4h19m
kube-public   Active   4h19m
kube-system   Active   4h19m

```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active   4h19m
kube-node-lease     Active   4h19m
kube-public         Active   4h19m
kube-system         Active   4h19m
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
aws-node-bmdht                     2/2     Running   0           4h14m
aws-node-qg58m                     2/2     Running   0           4h14m
coredns-86d7bdf-bhhxm             1/1     Running   0           4h16m
coredns-86d7bdf-j2wgl             1/1     Running   0           4h16m
kube-proxy-79g5s                   1/1     Running   0           4h14m
kube-proxy-qh6cf                   1/1     Running   0           4h14m
metrics-server-d9fd456dc-27jk9    1/1     Running   0           4h16m
metrics-server-d9fd456dc-zcwkg    1/1     Running   0           4h16m
ubuntu@ip-172-31-9-165:~$
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n kube-public
No resources found in kube-public namespace.
```

```
ubuntu@ip-172-31-9-165:~$ kubectl create ns my-namespace
namespace/my-namespace created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active   4h22m
kube-node-lease     Active   4h22m
kube-public         Active   4h22m
kube-system         Active   4h22m
my-namespace        Active   24s
```

```

ubuntu@ip-172-31-9-165:~$ kubectl get pods -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
aws-node-bmdht                      2/2     Running   0           4h14m
aws-node-qg58m                      2/2     Running   0           4h14m
coredns-86d7bdf-bhhxm               1/1     Running   0           4h16m
coredns-86d7bdf-j2wgl               1/1     Running   0           4h16m
kube-proxy-79g5s                    1/1     Running   0           4h14m
kube-proxy-qh6cf                    1/1     Running   0           4h14m
metrics-server-d9fd456dc-27jk9      1/1     Running   0           4h16m
metrics-server-d9fd456dc-zcwkg      1/1     Running   0           4h16m
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n kube-public
No resources found in kube-public namespace.
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl create ns my-namespace
namespace/my-namespace created
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active    4h22m
kube-node-lease     Active    4h22m
kube-public         Active    4h22m
kube-system         Active    4h22m
my-namespace        Active    24s

```

We can create a namespace in Kubernetes cluster: `kubectl create ns my-namespace`  
If needed, we can create multiple namespaces in Kubernetes cluster

To display all the available namespaces --> `$ kubectl get ns`

To get the pods available in specific namespace

`$ kubectl get pods -n <namespace-name>`

Ex: to get pods available in kube-system namespace

`$ kubectl get pods -n kube-system`

We can create a namespace in K8s cluster in 2 ways:

1. Using kubectl create ns command: `kubectl create ns my-namespace`
2. Using yaml file:

```
-rw-rw-r-- 1 ubuntu ubuntu 195 May 18 22:05 k8s-service-manifest.yml
```

```
ubuntu@ip-172-31-9-165:~$ vi k8s-namespace.yml
```

```
ubuntu@ip-172-31-9-165:~$ cat k8s-namespace.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: my-namespace
```

```
...
```

```
ubuntu@ip-172-31-9-165:~$ cat k8s-namespace.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

name: my-namespace

...

ubuntu@ip-172-31-9-165:~\$

ubuntu@ip-172-31-9-165:~\$ kubectl apply -f k8s-namespace.yml

Warning: resource namespaces/my-namespace is missing the kubectrl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.

namespace/my-namespace configured

ubuntu@ip-172-31-9-165:~\$ kubectl get ns

| NAME            | STATUS | AGE  |
|-----------------|--------|------|
| default         | Active | 5h1m |
| kube-node-lease | Active | 5h1m |
| kube-public     | Active | 5h1m |
| kube-system     | Active | 5h1m |
| my-namespace    | Active | 38m  |

ubuntu@ip-172-31-9-165:~\$ vi k8s-namespace.yml

ubuntu@ip-172-31-9-165:~\$ kubectl apply -f k8s-namespace.yml

namespace/my-namespace-2 created

ubuntu@ip-172-31-9-165:~\$ cat k8s-namespace.yml

---

apiVersion: v1

kind: Namespace

metadata:

name: my-namespace-2

...

ubuntu@ip-172-31-9-165:~\$ kubectl get ns

| NAME            | STATUS | AGE  |
|-----------------|--------|------|
| default         | Active | 5h3m |
| kube-node-lease | Active | 5h3m |
| kube-public     | Active | 5h3m |
| kube-system     | Active | 5h3m |
| my-namespace    | Active | 40m  |
| my-namespace-2  | Active | 15s  |

```

my-namespace-2 Active 40m
ubuntu@ip-172-31-9-165:~$ vi k8s-namespace.yml
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-namespace.yml
namespace/my-namespace-2 created
ubuntu@ip-172-31-9-165:~$ cat k8s-namespace.yml
---
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace-2
...
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS AGE
default             Active 5h3m
kube-node-lease     Active 5h3m
kube-public         Active 5h3m
kube-system         Active 5h3m
my-namespace        Active 40m
my-namespace-2      Active 15s
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl delete ns my-namespace
namespace "my-namespace" deleted

```

```

ubuntu@ip-172-31-9-165:~$ kubectl delete ns my-namespace
namespace "my-namespace" deleted
ubuntu@ip-172-31-9-165:~$ kubectl delete ns my-namespace-2
namespace "my-namespace-2" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS AGE
default             Active 5h19m
kube-node-lease     Active 5h19m
kube-public         Active 5h19m
kube-system         Active 5h19m

```

```

---
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace-1
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  namespace: my-namespace-1
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080

```

```

---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: LoadBalancer
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...
ubuntu@ip-172-31-9-165:~$ vi k8s-pod-svc-ns.yml
ubuntu@ip-172-31-9-165:~$ cat k8s-pod-svc-ns.yml
---
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace-1
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapp
  namespace: my-namespace-1
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: LoadBalancer
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...

We create another file due to indentation issues
ubuntu@ip-172-31-9-165:~$ cat k8s-pod-svc-ns1.yml
---
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace-1
---
apiVersion: v1

```

```
kind: Pod
metadata:
  name: javawebapp
  namespace: my-namespace-1
  labels:
    app: javawebapp
spec:
  containers:
    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
```

---

```
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
  namespace: my-namespace-1
spec:
  type: LoadBalancer
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
```

...

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-ns1.yml
namespace/my-namespace-1 unchanged
pod/javawebapp created
service/javawebappsvc created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
```

This time pods are getting created in a specific namespace and not default namespace

```
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME          STATUS  AGE
default       Active  5h46m
kube-node-lease Active  5h46m
kube-public   Active  5h46m
kube-system   Active  5h46m
my-namespace-1 Active  18m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME      READY  STATUS   RESTARTS  AGE
javawebapp 1/1    Running  0          10m
```



```

ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active   5h46m
kube-node-lease     Active   5h46m
kube-public         Active   5h46m
kube-system         Active   5h46m
my-namespace-1      Active   18m
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME                READY    STATUS    RESTARTS   AGE
javawebapp          1/1      Running   0           10m
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl get svc -n my-namespace-1
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)
AGE
javawebappsvc       LoadBalancer       10.100.191.133  a457deb47f843463b8f855e9ec0630e8-
250473956.ca-central-1.elb.amazonaws.com  80:30478/TCP  11m

```

```

ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active   5h46m
kube-node-lease     Active   5h46m
kube-public         Active   5h46m
kube-system         Active   5h46m
my-namespace-1      Active   18m
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME                READY    STATUS    RESTARTS   AGE
javawebapp          1/1      Running   0           10m
ubuntu@ip-172-31-9-165:~$ kubectl get svc -n my-namespace-1
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP
javawebappsvc       LoadBalancer       10.100.191.133  a457deb47f843463b8f855
ubuntu@ip-172-31-9-165:~$

```


A loadbalancer is created in EC2



## Load balancers (1/1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

🔍 Filter load balancers

| <input checked="" type="checkbox"/> | Name                   | DNS name                                                                                                    | State | VPC ID           |
|-------------------------------------|------------------------|-------------------------------------------------------------------------------------------------------------|-------|------------------|
| <input checked="" type="checkbox"/> | a457deb47f843463b8f... |  a457deb47f843463b8f855... | –     | vpc-0b9469b1bcf9 |

### Load balancer: a457deb47f843463b8f855e9ec0630e8

#### Details

##### Load balancer type

Classic

##### Status

2 of 2 instances in service

##### Scheme

Internet-facing

##### Hosted zone

ZQSVJUPU6J1EY

##### DNS name [Info](#)

 a457deb47f843463b8f855e9ec0630e8-250473956.ca-central-1.elb.amazonaws.com (A Record)


ubuntu@ip-172-31-9-165:~\$ kubectl get all -n my-namespace-1

| NAME           | READY | STATUS  | RESTARTS | AGE |
|----------------|-------|---------|----------|-----|
| pod/javawebapp | 1/1   | Running | 0        | 14m |

| NAME                  | TYPE         | CLUSTER-IP     | EXTERNAL-IP                                                               | PORT(S)      |
|-----------------------|--------------|----------------|---------------------------------------------------------------------------|--------------|
| service/javawebappsvc | LoadBalancer | 10.100.191.133 | a457deb47f843463b8f855e9ec0630e8-250473956.ca-central-1.elb.amazonaws.com | 80:30478/TCP |

Open the DNS url

<http://a457deb47f843463b8f855e9ec0630e8-250473956.ca-central-1.elb.amazonaws.com/>



Product Management System   View Products   Add Product   Update Product   Delete Product



```
ubuntu@ip-172-31-9-165:~$ kubectl delete ns my-namespace-1
namespace "my-namespace-1" deleted
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
No resources found in my-namespace-1 namespace.
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get ns
NAME                STATUS    AGE
default             Active    5h46m
kube-node-lease     Active    5h46m
kube-public         Active    5h46m
kube-system         Active    5h46m
my-namespace-1      Active    18m
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME                READY    STATUS    RESTARTS   AGE
javawebapp          1/1      Running   0           10m
ubuntu@ip-172-31-9-165:~$ kubectl get svc -n my-namespace-1
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP
javawebappsvc       LoadBalancer       10.100.191.133  a457deb47f843463b8f855e9ec0630e8-2
ubuntu@ip-172-31-9-165:~$ kubectl get all -n my-namespace-2
No resources found in my-namespace-2 namespace.
ubuntu@ip-172-31-9-165:~$ kubectl get all -n my-namespace-1
NAME                READY    STATUS    RESTARTS   AGE
pod/javawebapp      1/1      Running   0           14m

NAME                TYPE                CLUSTER-IP      EXTERNAL-IP
service/javawebappsvc LoadBalancer       10.100.191.133  a457deb47f843463b8f855e9ec
ubuntu@ip-172-31-9-165:~$ kubectl delete ns my-namespace-1
namespace "my-namespace-1" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
No resources found in my-namespace-1 namespace.
```

```
$ kubectl apply -f <yml>
$ kubectl get ns
$ kubectl get pods -n <namespace-name>
$ kubectl get service -n <namespace-name>
$ kubectl get all -n <namespace-name>
$ kubectl delete ns <namespace-name>
```

Container orchestration --> K8s introduction --> Advantages of K8s --> Architecture --> Components of architecture --> K8s cluster setup --> K8s resources --> POD, Services (ClusterIP, NodePort, LoadBalancer, Namespace)

Advantages of K8s:

- Self-healing (if any container/pod gets crashed, it will be replaced by new container)
- Load-balancing
- Auto-scaling

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-ns1.yml
namespace/my-namespace-1 created
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME                READY    STATUS    RESTARTS   AGE
javawebapp          1/1      Running   0           18s
```

```
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebapp -n my-namespace-1
pod "javawebapp" deleted
```

```

-rw-rw-r-- 1 ubuntu ubuntu 195 May 18 22:05 k8s-service-manifest.yml
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl apply -f k8s-pod-svc-ns1.yml
namespace/my-namespace-1 created
pod/javawebapp created
service/javawebappsvc created
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
NAME          READY   STATUS    RESTARTS   AGE
javawebapp    1/1     Running   0           18s
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebapp -n my-namespace-1
pod "javawebapp" deleted
ubuntu@ip-172-31-9-165:~$ █

```

Now I have deleted the pod.

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -n my-namespace-1
```

No resources found in my-namespace-1 namespace.

Is it self-healing? No pods are there. No

So far we have created POD directly using POD manifest YML (kind: Pod)

If we create POD directly then we don't get self-healing capability

If POD is damaged/crashed/deleted, then K8s will not create a new POD

In this case, if POD is damaged then our application will be down

Hence we should not create POD directly to deploy our application in K8s and we need to use the concept of K8s resources to create the PODs

If we create POD using K8s resources then Pod lifecycle will be managed by K8s

We have following resources to create PODs ->

1. ReplicationController (outdated)
2. ReplicaSet
3. Deployment
4. DaemonSet
5. StatefulSet

We have to use these methods to obtain the advantages or benefits of K8s

ReplicaSet:

It is one of the K8s resource used to create and manage Pods and ReplicaSet will take care of Pod lifecycle

When Pod is damaged/crashed/deleted, then ReplicaSet will create a new Pod

It will always maintain given number of Pod count for our application

For example, in the manifest file, if we say replicas: 2, it will maintain 2

It will help us to achieve high availability for our application. We can also scale-up and scale-down our pod count

replicaSet is a part of apps group

replicas: 2 I want 2 pods running at all the time

We are giving only 2 spacings

matchLabels: replicas for which pod

replicas : 2

```
$ vi replicaSet.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: ReplicaSet
```

```
metadata:
  name: javawebrs
spec:
  replicas: 2
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebpod
    labels:
      app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          ports:
            - containerPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~$ cat replicaSet.yml
```

```
---
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: javawebrs
spec:
  replicas: 2
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebpod
    labels:
      app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          ports:
            - containerPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f replicaSet.yml
replicaset.apps/javawebrs created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY STATUS RESTARTS AGE
pod/javawebrs-6cfpx 1/1 Running 0 11m
pod/javawebrs-76wx9 1/1 Running 0 11m
```

```
NAME          TYPE      CLUSTER-IP  EXTERNAL-IP  PORT(S)  AGE
service/kubernetes ClusterIP  10.100.0.1  <none>       443/TCP  5h22m
```

```
NAME          DESIRED  CURRENT  READY  AGE
```

```
replicaset.apps/javawebrs 2 2 2 11m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                READY   STATUS    RESTARTS   AGE
pod/javawebrs-6cfpx 1/1     Running   0           11m
pod/javawebrs-76wx9 1/1     Running   0           11m

NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes  ClusterIP     10.100.0.1   <none>        443/TCP    5h22m

NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/javawebrs 2         2         2       11m
ubuntu@ip-172-31-9-165:~$
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
javawebrs-6cfpx    1/1     Running   0           12m
javawebrs-76wx9    1/1     Running   0           12m
```

2 pods are there, we specified replicas as 2

```
ubuntu@ip-172-31-9-165:~$ kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
javawebrs          2         2         2       13m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                READY   STATUS    RESTARTS   AGE
pod/javawebrs-6cfpx 1/1     Running   0           11m
pod/javawebrs-76wx9 1/1     Running   0           11m
```

```
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes  ClusterIP     10.100.0.1   <none>        443/TCP    5h22m
```

```
NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/javawebrs 2         2         2       11m
```

```
ubuntu@ip-172-31-9-165:~$
```

```
ubuntu@ip-172-31-9-165:~$
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
javawebrs-6cfpx    1/1     Running   0           12m
javawebrs-76wx9    1/1     Running   0           12m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
javawebrs          2         2         2       13m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebrs-6cfpx
pod "javawebrs-6cfpx" deleted
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
javawebrs-76wx9    1/1     Running   0           14m
javawebrs-rb2kh    1/1     Running   0           24s
```

We deleted pod "javawebrs-6cfpx" but pod "javawebrs-76wx9" is created. So Self-healing



```

ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY    STATUS    RESTARTS   AGE
pod/javawebrs-6cfpx                 1/1     Running   0           11m
pod/javawebrs-76wx9                 1/1     Running   0           11m

NAME                                TYPE     CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP  10.100.0.1    <none>       443/TCP    5h22m

NAME                                DESIRED   CURRENT   READY    AGE
replicaset.apps/javawebrs           2         2         2        11m
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
javawebrs-6cfpx                     1/1     Running   0           12m
javawebrs-76wx9                     1/1     Running   0           12m
ubuntu@ip-172-31-9-165:~$ kubectl get rs
NAME                                DESIRED   CURRENT   READY    AGE
javawebrs                           2         2         2        13m
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebrs-6cfpx
pod "javawebrs-6cfpx" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
javawebrs-76wx9                     1/1     Running   0           14m
javawebrs-rb2kh                     1/1     Running   0           24s
ubuntu@ip-172-31-9-165:~$

```

```

javawebrs                           2         2         2        13m
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebrs-6cfpx
pod "javawebrs-6cfpx" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
javawebrs-76wx9                     1/1     Running   0           14m
javawebrs-rb2kh                     1/1     Running   0           24s
ubuntu@ip-172-31-9-165:~$ kubectl delete pod javawebrs-76wx9
pod "javawebrs-76wx9" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
javawebrs-mv6qd                     1/1     Running   0           5s
javawebrs-rb2kh                     1/1     Running   0           7m43s
ubuntu@ip-172-31-9-165:~$

```

On run I am changing replicaSet.yml: replicas to 4

---

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: javawebrs

spec:

replicas: 4

selector:

matchLabels:

app: javawebapp

template:

metadata:

name: javawebpod

labels:

app: javawebapp

spec:

containers:

```
- name: javawebappcontainer
  image: hacker123shiva/springbt-in-docker:latest
  ports:
    - containerPort: 8080
...

```

Not newly created it is only configured

```
ubuntu@ip-172-31-9-165:~$ vi replicaSet.yml
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f replicaSet.yml
```

replicaset.apps/javawebrs configured

Now we have 4 pods

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
```

| NAME            | READY | STATUS  | RESTARTS | AGE |
|-----------------|-------|---------|----------|-----|
| javawebrs-k7tgg | 1/1   | Running | 0        | 29s |
| javawebrs-mv6qd | 1/1   | Running | 0        | 12m |
| javawebrs-rb2kh | 1/1   | Running | 0        | 20m |
| javawebrs-rg9g8 | 1/1   | Running | 0        | 29s |

```
ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 6
```

```
ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 6
```

replicaset.apps/javawebrs scaled

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods
```

| NAME            | READY | STATUS  | RESTARTS | AGE |
|-----------------|-------|---------|----------|-----|
| javawebrs-b6ffd | 1/1   | Running | 0        | 6s  |
| javawebrs-k7tgg | 1/1   | Running | 0        | 2m  |
| javawebrs-mv6qd | 1/1   | Running | 0        | 14m |
| javawebrs-mzk67 | 1/1   | Running | 0        | 6s  |
| javawebrs-rb2kh | 1/1   | Running | 0        | 21m |
| javawebrs-rg9g8 | 1/1   | Running | 0        | 2m  |

```
javawebrs-rb2kh 1/1 Running 0 7m43s
ubuntu@ip-172-31-9-165:~$ vi replicaSet.yml
ubuntu@ip-172-31-9-165:~$ kubectl apply -f replicaSet.yml
replicaset.apps/javawebrs configured
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
javawebrs-k7tgg 1/1     Running   0          29s
javawebrs-mv6qd 1/1     Running   0          12m
javawebrs-rb2kh 1/1     Running   0          20m
javawebrs-rg9g8 1/1     Running   0          29s
ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 6
replicaset.apps/javawebrs scaled
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
javawebrs-b6ffd 1/1     Running   0          6s
javawebrs-k7tgg 1/1     Running   0          2m
javawebrs-mv6qd 1/1     Running   0          14m
javawebrs-mzk67 1/1     Running   0          6s
javawebrs-rb2kh 1/1     Running   0          21m
javawebrs-rg9g8 1/1     Running   0          2m
ubuntu@ip-172-31-9-165:~$
```

Now scaled down to 2

```

ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 2
replicaset.apps/javawebrs scaled
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY STATUS  RESTARTS  AGE
javawebrs-mv6qd 1/1   Running  0         15m
javawebrs-rb2kh 1/1   Running  0         23m
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 6
replicaset.apps/javawebrs scaled
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY STATUS  RESTARTS  AGE
javawebrs-b6ffd 1/1   Running  0         6s
javawebrs-k7tgg 1/1   Running  0         2m
javawebrs-mv6qd 1/1   Running  0         14m
javawebrs-mzk67 1/1   Running  0         6s
javawebrs-rb2kh 1/1   Running  0         21m
javawebrs-rg9g8 1/1   Running  0         2m
ubuntu@ip-172-31-9-165:~$ kubectl scale rs javawebrs --replicas 2
replicaset.apps/javawebrs scaled
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME          READY STATUS  RESTARTS  AGE
javawebrs-mv6qd 1/1   Running  0         15m
javawebrs-rb2kh 1/1   Running  0         23m
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl delete rs javawebrs
replicaset.apps "javawebrs" deleted
ubuntu@ip-172-31-9-165:~$ kubectl get pods
No resources found in default namespace.

```

If we want to delete the pods, then we have to delete the resource, which created the pods  
\$ kubectl delete rs javawebrs

In ReplicaSet, scale up and scale down is a manual process  
K8s supports auto-scaling when we use "Deployment" resource to create Pods

#### K8s Deployment

It is one of the K8s resource/component recommended to deploy our application where Deployment will manage Pod lifecycle

Advantages with K8s deployment

1. Zero downtime
2. Auto scaling
3. Rolling update (it will delete and create new pods one by one) and Rollback

When I deploy application, existing pods (old pods will have old application) should be replaced by new pods with application (new pods will have new application). Rolling update means it will delete one of the old pods then create a new pod, same delete second old one replace with a new pod

We have deployment strategies:

1. Recreate (delete all existing pods and create new pods)
2. Rolling update (it will delete and create new pod one by one)

```

ubuntu@ip-172-31-9-165:~$ vi deployment.yml
ubuntu@ip-172-31-9-165:~$ cat deployment.yml
---

```



```

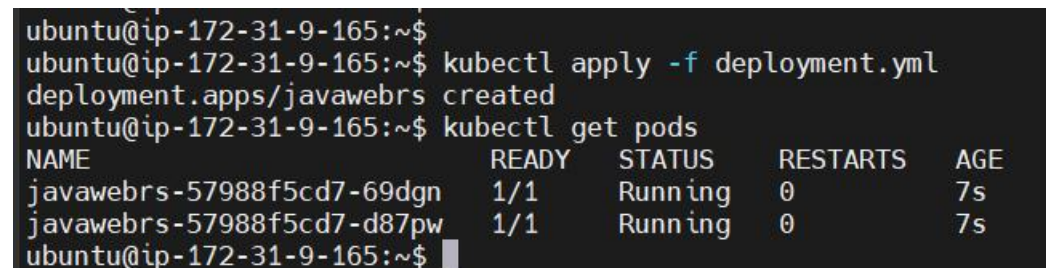
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebrs
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebpod
    labels:
      app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          ports:
            - containerPort: 8080
...

```

```

ubuntu@ip-172-31-9-165:~$ kubectl apply -f deployment.yml
deployment.apps/javawebrs created
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY  STATUS   RESTARTS  AGE
javawebrs-57988f5cd7-69dgn         1/1    Running  0          7s
javawebrs-57988f5cd7-d87pw         1/1    Running  0          7s

```



```

ubuntu@ip-172-31-9-165:~$
ubuntu@ip-172-31-9-165:~$ kubectl apply -f deployment.yml
deployment.apps/javawebrs created
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY  STATUS   RESTARTS  AGE
javawebrs-57988f5cd7-69dgn         1/1    Running  0          7s
javawebrs-57988f5cd7-d87pw         1/1    Running  0          7s
ubuntu@ip-172-31-9-165:~$

```

```

ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY  STATUS   RESTARTS  AGE
pod/javawebrs-57988f5cd7-69dgn     1/1    Running  0          45s
pod/javawebrs-57988f5cd7-d87pw     1/1    Running  0          45s

```

```

NAME      TYPE      CLUSTER-IP  EXTERNAL-IP  PORT(S)  AGE
service/kubernetes  ClusterIP  10.100.0.1  <none>       443/TCP  7h9m

```

```

NAME                                READY  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/javawebrs           2/2    2           2          45s

```

```

NAME                                DESIRED  CURRENT  READY  AGE
replicaset.apps/javawebrs-57988f5cd7  2        2        2      45s

```

```

javawebbrs-57988f5cd7-d87pw 1/1 Running 0 45s
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/javawebbrs-57988f5cd7-69dgn    1/1    Running   0           45s
pod/javawebbrs-57988f5cd7-d87pw    1/1    Running   0           45s

NAME                                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes                  ClusterIP     10.100.0.1   <none>        443/TCP    7h9m

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/javawebbrs          2/2     2             2           45s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/javawebbrs-57988f5cd7 2         2         2       45s
ubuntu@ip-172-31-9-165:~$

```

ubuntu@ip-172-31-9-165:~\$ kubectl get deployment

```

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
javawebbrs 2/2     2           2           92s

```

Updated yml file

ubuntu@ip-172-31-9-165:~\$ cat deployment.yml

---

apiVersion: apps/v1

kind: Deployment

metadata:

name: javawebdeployment

spec:

replicas: 2

strategy:

type: RollingUpdate

selector:

matchLabels:

app: javawebapp

template:

metadata:

name: javawebpod

labels:

app: javawebapp

spec:

containers:

- name: javawebappcontainer

image: hacker123shiva/springbt-in-docker:latest

ports:

- containerPort: 8080

...

ubuntu@ip-172-31-9-165:~\$ kubectl apply -f deployment.yml

deployment.apps/javawebdeployment created

ubuntu@ip-172-31-9-165:~\$ kubectl get pods

```

NAME                                READY   STATUS    RESTARTS   AGE
javawebdeployment-57988f5cd7-6gnjl 1/1    Running   0           4s
javawebdeployment-57988f5cd7-c4dxc 1/1    Running   0           4s

```

ubuntu@ip-172-31-9-165:~\$ kubectl get deployment

```

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
javawebdeployment 2/2     2           2           12s

```

```
ubuntu@ip-172-31-9-165:~$ kubectl scale deployment javawebdeployment --replicas 4
deployment.apps/javawebdeployment scaled
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
```

| NAME                                   | READY | STATUS  | RESTARTS | AGE |
|----------------------------------------|-------|---------|----------|-----|
| pod/javawebdeployment-57988f5cd7-2ktmv | 1/1   | Running | 0        | 20s |
| pod/javawebdeployment-57988f5cd7-6gnjl | 1/1   | Running | 0        | 65s |
| pod/javawebdeployment-57988f5cd7-98ggd | 1/1   | Running | 0        | 20s |
| pod/javawebdeployment-57988f5cd7-c4dxc | 1/1   | Running | 0        | 65s |

| NAME               | TYPE      | CLUSTER-IP | EXTERNAL-IP | PORT(S) | AGE   |
|--------------------|-----------|------------|-------------|---------|-------|
| service/kubernetes | ClusterIP | 10.100.0.1 | <none>      | 443/TCP | 7h12m |

| NAME                              | READY | UP-TO-DATE | AVAILABLE | AGE |
|-----------------------------------|-------|------------|-----------|-----|
| deployment.apps/javawebdeployment | 4/4   | 4          | 4         | 65s |

| NAME                                         | DESIRED | CURRENT | READY | AGE |
|----------------------------------------------|---------|---------|-------|-----|
| replicaset.apps/javawebdeployment-57988f5cd7 | 4       | 4       | 4     | 65s |

Auto-scaling can be done in two ways: Horizontal scaling and Vertical scaling

To do more work, I have two options, give more pay to existing engineers (Vertical) or hire new engineers (Horizontal)

Vertical scaling, existing engineers should do more work and leads to less productivity so Horizontal scaling is better

Auto-scaling focuses on Horizontal scaling

Horizontal Pod Autoscaling (HPA): it is used to scale up and down number of pods/replicas based on observed metrics (CPU/memory utilization). To achieve autoscaling, we first have to create metric server. HPA will interact with this metric server to identify CPU/memory utilization, based on that Horizontal Autoscaling will be done. Creating new machines is Horizontal

```
kubectl delete all --all
eksctl delete cluster --name my-eks-cluster --region ca-central-1
```