

Kubernetes_4

```
eksctl create cluster --name my-eks-cluster --region ca-central-1 --node-type t2.medium --zones ca-central-1a,ca-central-1b
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f deployment.yml
deployment.apps/javawebdeployment created
ubuntu@ip-172-31-9-165:~$
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY STATUS RESTARTS AGE
pod/javawebdeployment-57988f5cd7-9g8tr 1/1 Running 0      21s
pod/javawebdeployment-57988f5cd7-p7cb7 1/1 Running 0      22s
```

```
NAME          TYPE      CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/kubernetes ClusterIP 10.100.0.1 <none>    443/TCP 8m11s
```

```
NAME                                READY UP-TO-DATE AVAILABLE AGE
deployment.apps/javawebdeployment 2/2   2         2      22s
```

```
NAME                                DESIRED CURRENT READY AGE
replicaset.apps/javawebdeployment-57988f5cd7 2      2      2      22s
```

```
ubuntu@ip-172-31-9-165:~$ kubectl scale deployment javawebdeployment --replicas 4
deployment.apps/javawebdeployment scaled
ubuntu@ip-172-31-9-165:~$ kubectl get pods
NAME                                READY STATUS RESTARTS AGE
javawebdeployment-57988f5cd7-9g8tr 1/1 Running 0      2m
javawebdeployment-57988f5cd7-g2th6 1/1 Running 0      5s
javawebdeployment-57988f5cd7-p7cb7 1/1 Running 0      2m1s
javawebdeployment-57988f5cd7-t6xpg 1/1 Running 0      5s
ubuntu@ip-172-31-9-165:~$
```

If you don't explicitly specify service, the default service is ClusterIP

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebdeploy
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebpod
    labels:
      app: javawebapp
    spec:
      containers:
```

```

    - name: javawebappcontainer
      image: hacker123shiva/springbt-in-docker:latest
      ports:
        - containerPort: 8080
  ---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: LoadBalancer
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...

```

ubuntu@ip-172-31-9-165:~\$ cat dep-svc.yml

```

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebdeployment
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebpod
      labels:
        app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          ports:
            - containerPort: 8080
  ---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: LoadBalancer
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
...

```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f dep-svc.yml
deployment.apps/javawebdeployment created
service/javawebappsvc created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get all
NAME                                READY STATUS RESTARTS AGE
pod/javawebdeployment-57988f5cd7-4mhbd 1/1   Running 0      70s
pod/javawebdeployment-57988f5cd7-wm8gz 1/1   Running 0      70s
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/javawebappsvc	LoadBalancer	10.100.177.181	a2c8cab06e2fa4241aec665485af8c3c-58365812.ca-central-1.elb.amazonaws.com	80:31546/TCP	70s
service/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	7m48s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/javawebdeployment	2/2	2	2	71s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/javawebdeployment-57988f5cd7	2	2	2	71s

Go to EC2 Loadbalancer

Load balancers (1/1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

<input checked="" type="checkbox"/>	Name	DNS name	State	VPC ID
<input checked="" type="checkbox"/>	a2c8cab06e2fa4241aec...	a2c8cab06e2fa4241aec665...	-	vpc-04028b965c79a1c2f

Load balancer: a2c8cab06e2fa4241aec665485af8c3c

Details

Listeners

Network mapping

Security

Health checks

Target instances

Details

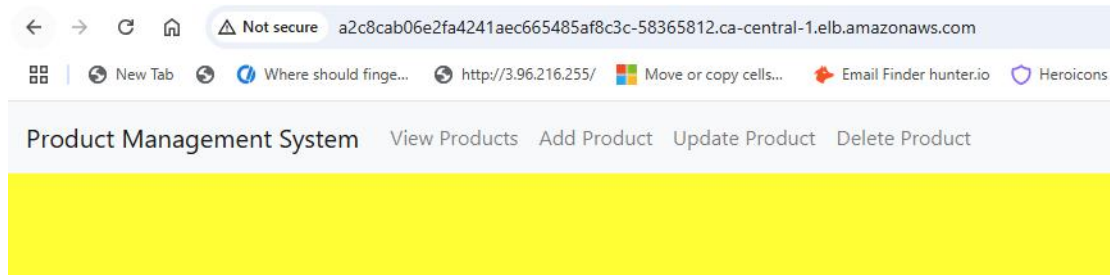
Load balancer type
Classic

Scheme
Internet-facing

Status
2 of 2 instances in service

Hosted zone
ZQSVJUPU6J1EY

a2c8cab06e2fa4241aec665485af8c3c-58365812.ca-central-1.elb.amazonaws.com



```
ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "javawebdeployment-57988f5cd7-4mhbd" deleted
pod "javawebdeployment-57988f5cd7-wm8gz" deleted
service "javawebappsvc" deleted
service "kubernetes" deleted
deployment.apps "javawebdeployment" deleted
replicaset.apps "javawebdeployment-57988f5cd7" deleted
```

Autoscaling

Process of increasing or decreasing the infrastructure resources based on the demand

Autoscaling: Horizontal (increasing / decreasing infrastructure resources based on the demand) and Vertical scaling (increasing the capacity of the same single system / machine / pod)

HPA (Horizontal Pod Autoscaling) and VPA (Vertical Pod Autoscaling)

Why Autoscale in K8s?

High availability of application or better availability of application

Elasticity

Efficient resource utilization

Seamless load management

HPA (Horizontal Pod Autoscaling) --> used to scale up or scale down number of pod replicas based on observed metrics (CPU or memory utilization). we cannot simply add or remove resources, we got to check certain metrics before we hire someone

HPA needs metrics to adjust the pods

HPA observes all metrics --> based on the observation, it will add/delete pods, tracks multiple metrics

HPA will interact with Metric server to identify CPU/memory utilization of POD

Metric server is an application that collects metrics from objects, pods, nodes according to state of CPU and memory

```
ubuntu@ip-172-31-9-165:~$ kubectl top nodes
```

NAME	CPU(cores)	CPU(%)	MEMORY(bytes)	MEMORY(%)
ip-192-168-13-3.ca-central-1.compute.internal	39m	2%	629Mi	18%
ip-192-168-34-135.ca-central-1.compute.internal	24m	1%	550Mi	16%

Metrics server will not be present by default in Kubernetes server

Install metrics API

```
$ mkdir k8s-metrics-server
```

```
$ cd k8s-metrics-server
```

```
$ vi metrics-api-server.yaml
```

```
---
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: metrics
```

```
...
```

```
ubuntu@ip-172-31-9-165:~$ cat metrics-api-server.yaml
```

```
---
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: metrics-server
  name: metrics-server
  namespace: metrics
```

```
...
```

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ cat metrics-deployment.yaml
```

```
---
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: metrics-server
  namespace: kube-system
  labels:
    k8s-app: metrics-server
spec:
  selector:
    matchLabels:
      k8s-app: metrics-server
      app.kubernetes.io/instance: metrics-server
      app.kubernetes.io/name: metrics-server
  strategy:
    rollingUpdate:
      maxUnavailable: 0
  template:
    metadata:
      labels:
        k8s-app: metrics-server
        app.kubernetes.io/instance: metrics-server
        app.kubernetes.io/name: metrics-server
    spec:
      serviceAccountName: metrics-server
      nodeSelector:
        kubernetes.io/os: linux
      priorityClassName: system-cluster-critical
      containers:
        - name: metrics-server
          image: k8s.gcr.io/metrics-server/metrics-server:v0.5.0
          imagePullPolicy: IfNotPresent
          args:
            - --cert-dir=/tmp
```

```

    --secure-port=443
    --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
    --kubelet-use-node-status-port
    --metric-resolution=15s
    --kubelet-insecure-tls
  ports:
    - containerPort: 443
      name: https
      protocol: TCP
  livenessProbe:
    failureThreshold: 3
    httpGet:
      path: /livez
      port: https
      scheme: HTTPS
    periodSeconds: 10
  readinessProbe:
    failureThreshold: 3
    httpGet:
      path: /readyz
      port: https
      scheme: HTTPS
    initialDelaySeconds: 20
    periodSeconds: 10
  resources:
    requests:
      cpu: 100m
      memory: 200Mi
  securityContext:
    readOnlyRootFilesystem: true
    runAsNonRoot: true
    runAsUser: 1000
  volumeMounts:
    - mountPath: /tmp
      name: tmp-dir
  volumes:
    - name: tmp-dir
      emptyDir: {}
...

```

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ cat metrics-rbac.yaml
```

```

---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: system:aggregated-metrics-reader
rules:
  - apiGroups: ["metrics.k8s.io"]
    resources: ["pods", "nodes"]
    verbs: ["get", "list", "watch"]

```

```

---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: system:metrics-server
rules:
- apiGroups: ["" ]
  resources:
    - pods
    - nodes
    - nodes/stats
    - namespaces
    - configmaps
    - services
  verbs: ["get", "list", "watch"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: metrics-server:system:auth-delegator
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:auth-delegator
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: system:metrics-server
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:metrics-server
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system
...

```

ubuntu@ip-172-31-9-165:~/k8s-metrics-server\$ cat metrics-server-service.yaml

```

---
apiVersion: v1
kind: Service
metadata:
  name: metrics-server
  namespace: kube-system
labels:
  k8s-app: metrics-server
spec:
  selector:
    k8s-app: metrics-server

```

```
ports:
- port: 443
  targetPort: https
  protocol: TCP
  name: https
type: ClusterIP
...
```

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ cat metrics-serviceaccount.yaml
```

```
---
apiVersion: apiregistration.k8s.io/v1
kind: APIService
metadata:
  name: v1beta1.metrics.k8s.io
  labels:
    k8s-app: metrics-server
spec:
  group: metrics.k8s.io
  version: v1beta1
  insecureSkipTLSVerify: true
  groupPriorityMinimum: 100
  versionPriority: 100
  service:
    name: metrics-server
    namespace: kube-system
...
```

Create a metric system inside the namespace: kube-system

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ ls -l
total 20
-rw-rw-r-- 1 ubuntu ubuntu 138 May 25 18:42 metrics-api-server.yaml
-rw-rw-r-- 1 ubuntu ubuntu 923 May 25 18:44 metrics-deployment.yaml
-rw-rw-r-- 1 ubuntu ubuntu 1212 May 25 18:48 metrics-rbac.yaml
-rw-rw-r-- 1 ubuntu ubuntu 288 May 25 18:52 metrics-server-service.yaml
-rw-rw-r-- 1 ubuntu ubuntu 102 May 25 18:55 metrics-serviceaccount.yaml
```

Run or Execute the Yaml file

\$ kubectl apply -f k8s-metrics-server (run directory only if all yaml files inside the dir are correct)
Recommended approach: Run all yaml files individually

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ kubectl apply -f metrics-api-server.yaml
Warning: resource serviceaccounts/metrics-server is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.
serviceaccount/metrics-server configured
```

```
kubectl apply -f metrics-deployment.yaml
deployment.apps/metrics-server created
```

```
ubuntu@ip-172-31-9-165:~/k8s-metrics-server$ kubectl apply -f metrics-rbac.yaml
serviceaccount/metrics-server configured
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
```


Warning: resource clusterroles/system:metrics-server is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

clusterrole.rbac.authorization.k8s.io/system:metrics-server configured

Warning: resource clusterrolebindings/metrics-server:system:auth-delegator is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator configured

Warning: resource clusterrolebindings/system:metrics-server is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server configured

ubuntu@ip-172-31-9-165:~/k8s-metrics-server\$ kubect apply -f metrics-server-service.yaml

Warning: resource services/metrics-server is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

service/metrics-server configured

ubuntu@ip-172-31-9-165:~/k8s-metrics-server\$ kubect apply -f metrics-serviceaccount.yaml

Warning: resource apiservices/v1beta1.metrics.k8s.io is missing the kubect.kubernetes.io/last-applied-configuration annotation which is required by kubect apply. kubect apply should only be used on resources created declaratively by either kubect create --save-config or kubect apply. The missing annotation will be patched automatically.

apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io configured

It will create metrics-server, service-account, role, role binding and all the stuff required

Check top nodes using metrics-server

ubuntu@ip-172-31-9-165:~\$ eksctl create addon --name metrics-server --cluster my-eks-cluster --region ca-central-1 --force

ubuntu@ip-172-31-9-165:~\$ kubect get pods -n kube-system -l k8s-app=metrics-server

NAME	READY	STATUS	RESTARTS	AGE
metrics-server-79bb88c6fc-gv6mv	0/1	Running	0	11m

kubect -n kube-system exec -it metrics-server-79bb88c6fc-gv6mv -- /bin/sh
curl -k https://<worker-node-ip>:10250/stats/summary

I will debug later

ubuntu@ip-172-31-9-165:~\$ kubect top nodes

error: Metrics API not available

<https://docs.aws.amazon.com/eks/latest/userguide/metrics-server.html>

This is working fine

kubect apply -f <https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>

kubect get deployment metrics-server -n kube-system

kubectl top nodes

```
ubuntu@ip-172-31-9-165:~$ kubectl get deployment metrics-server -n kube-system
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
metrics-server	1/1	1	1	84s

```
ubuntu@ip-172-31-9-165:~$ kubectl top nodes
```

NAME	CPU(cores)	CPU(%)	MEMORY(bytes)	MEMORY(%)
ip-192-168-13-3.ca-central-1.compute.internal	28m	1%	622Mi	18%
ip-192-168-34-135.ca-central-1.compute.internal	28m	1%	617Mi	18%

```
ubuntu@ip-172-31-9-165:~$ kubectl top pods
```

No resources found in default namespace.

Note: metrics-server will be installed under kube-system namespace

Deploy same application

```
ubuntu@ip-172-31-9-165:~$ cat hpa-demo-deployment.yaml
```

apiVersion: apps/v1

kind: Deployment

metadata:

name: hpa-demo-deployment

labels:

app: hpa-demo

spec:

replicas: 1

selector:

matchLabels:

app: hpa-demo

template:

metadata:

labels:

app: hpa-demo

spec:

containers:

- name: hpa-demo-container

image: k8s.gcr.io/hpa-example

ports:

- containerPort: 80

resources:

requests:

cpu: 100m

limits:

cpu: 500m

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f hpa-demo-deployment.yaml
```

deployment.apps/hpa-demo-deployment created

Create HPA service

```
ubuntu@ip-172-31-9-165:~$ cat hpa-demo-service.yaml
```

apiVersion: v1

kind: Service

metadata:

name: hpa-demo-service

labels:

run: hpa-demo

```
spec:
  selector:
    run: hpa-demo
  ports:
  - port: 80
    targetPort: 80
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f hpa-demo-service.yaml
service/hpa-demo-service created
```

Create HPA

```
ubuntu@ip-172-31-9-165:~$ cat hpa-demo.yaml
```

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hpa-demo-hpa
  namespace: default
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hpa-demo-deployment
  minReplicas: 1
  maxReplicas: 5
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```

```
ubuntu@ip-172-31-9-165:~$ kubectl apply -f hpa-demo.yaml
horizontalpodautoscaler.autoscaling/hpa-demo-hpa created
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get deploy
NAME                READY  UP-TO-DATE  AVAILABLE  AGE
hpa-demo-deployment 1/1    1           1           25m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get svc
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
hpa-demo-service    ClusterIP   10.100.42.151 <none>       80/TCP     18m
kubernetes          ClusterIP   10.100.0.1    <none>       443/TCP    5h57m
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa
NAME                REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
hpa-demo-hpa        Deployment/hpa-demo-deployment  cpu: 1%/50%  1        5        1        8m39s
```

To demonstrate auto-scaling, I will increase load on this machine

```
kubectl run -i --tty load-generator --image=busybox --restart=Never -- /bin/sh -c "while sleep 0.01; do
wget -q -o- http://hpa-demo; done"
```

```
ubuntu@ip-172-31-9-165:~$ kubectl get pods -l app=hpa-demo
```

NAME	READY	STATUS	RESTARTS	AGE
hpa-demo-deployment-7577d65cb7-ckp7c	1/1	Running	0	48m

OK!OK!OK!OK!OK!^Cpod default/load-generator terminated (Error)

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 44%/50%	1	5	5	43m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 119%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	44m

As load increases, POD increases, reached max 5

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa -w
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 44%/50%	1	5	5	43m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 119%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 131%/50%	1	5	5	44m

I stopped the load balancer

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa -w
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 44%/50%	1	5	5	43m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 119%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 131%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 123%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 53%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 1%/50%	1	5	5	46m

Now the load has dropped

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa -w
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 44%/50%	1	5	5	43m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 119%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 131%/50%	1	5	5	44m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 123%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 125%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 53%/50%	1	5	5	45m
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 1%/50%	1	5	5	46m

After sometime, replicas is scaled-down as well

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 1%/50%	1	5	1	52m

```
ubuntu@ip-172-31-9-165:~$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
hpa-demo-hpa	Deployment/hpa-demo-deployment	cpu: 1%/50%	1	5	1	52m

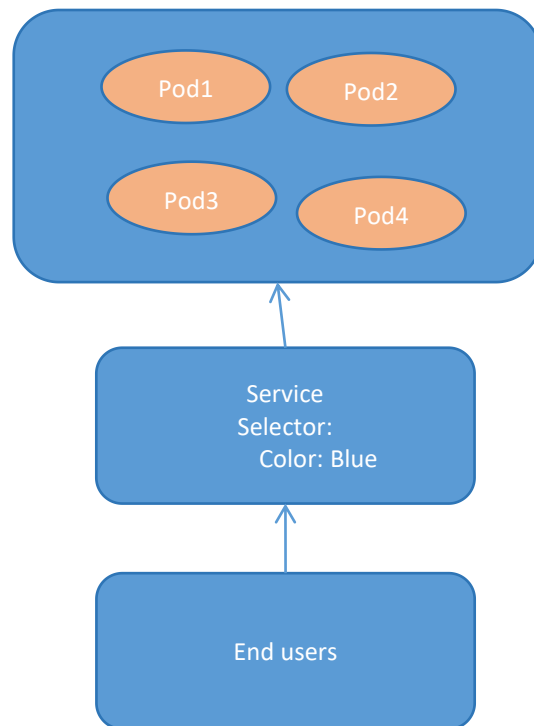
```
ubuntu@ip-172-31-9-165:~$
```

Container orchestration --> K8s introduction --> Advantages --> Architecture --> Architecture components --> K8s cluster setup --> K8s resources --> POD, Service (ClusterIP, NodePort & LoadBalancer) --> Namespace --> ReplicaSet --> Deployment --> Metrics-server --> HPA

Blue-Green deployment model -> strategy of application deployment

Assume that application is running on different PODs. We cannot access the Pods directly so we create a Service to expose them. Endusers are accessing our application through Service. Application is inside the Pods, and after 10 days there are major updates that we want to do. If my Endusers are using the same Pods, inside Pods, application will be there. If we are exposing the same old Pods, then Endusers are using only the same Old applications. Every application changes over time. I want to have new Pods, if new Pods need to work, then it is a Deployment process. Deployment takes time and there will be some downtime. If I delete my old pods and create new pods, by default application downtime will be there. What if you have some errors or issues in your latest Pods. Don't you think it is a big business loss. That's where the concept of Blue-Green deployment comes into picture. It is like an application-release model, it will help you to decrease the downtime. It will follow a strategy in a

way that downtime will also be reduced and if there is new problem or issue with the new Pods, Endusers can rollback / go back to the Old Pods with minimal time. It will show you how to release your application with minimal downtime. It is decreasing the risks we face during application release



If I say Selector: Color: Blue, people are able to access it. Then I will create another deployment and I will label as Green. I have created new Pods and this is deployed recently. Now before I give access to Endusers to access the new Pods, I will expose only to our testing team with the Pre-Production Service. Pre-Prod service is accessing the application only for the testing purpose. Inside the new Service, I will give the Selector as Green. Even if there is issues after new pods are deployed, our Endusers are not getting affected. Our testing team is now ok with the new pods. Then I change Selector color of Blue to Green. This is not re-deployment, it is not taking much time. New pods are already created, up and running, so no downtime. Just trying to change the Service from Blue to Green it won't take much time. We could add a Router/Live service also that would re-direct Users to Blue or Green environment. If Green environment is stable then we rename as Blue then new Green environment will start.

Blue-Green deployment is an application release model, which decreases risk and minimizes downtime. It uses two production environments known as Blue and Green. Old version of the application is called as the Blue environment and new version is known as Green environment. Four Yaml files are required: Blue deployment, Green deployment, Live service, Pre-prod service

```

ubuntu@ip-172-31-9-165:~$ kubectl delete deployment metrics-server -n kube-system
deployment.apps "metrics-server" deleted
ubuntu@ip-172-31-9-165:~$ kubectl delete all --all
pod "hpa-demo-deployment-7577d65cb7-8tdh7" deleted
pod "load-generator" deleted
service "hpa-demo-service" deleted
service "kubernetes" deleted
deployment.apps "hpa-demo-deployment" deleted
horizontalpodautoscaler.autoscaling "hpa-demo-hpa" deleted
  
```

```
buntu@ip-172-31-9-165:~$ mkdir blue-green-model
ubuntu@ip-172-31-9-165:~$ cd blue-green-model/
ubuntu@ip-172-31-9-165:~/blue-green-model$ ls -l
total 0
ubuntu@ip-172-31-9-165:~/blue-green-model$ touch blue-deployment.yml
ubuntu@ip-172-31-9-165:~/blue-green-model$ touch live-service.yml
ubuntu@ip-172-31-9-165:~/blue-green-model$ touch green-deployment.yml
ubuntu@ip-172-31-9-165:~/blue-green-model$ touch pre-pod.yml
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat blue-deployment.yml
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebbbluedeploy
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: java-web-app
      version: v1
      color: blue
  template:
    metadata:
      labels:
        app: java-web-app
        version: v1
        color: blue
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat live-service.yml
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: javawebapplivesvc
spec:
  type: LoadBalancer
  selector:
    app: java-web-app # Matches the app
    color: blue # Sends traffic to the blue pods
  ports:
    - port: 80
      targetPort: 8080
...
```



```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f blue-deployment.yml
deployment.apps/javawebbbluedeploy created
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f live-service.yml
service/javawebbapplivesvc created
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get pods
NAME                                READY  STATUS   RESTARTS  AGE
javawebbbluedeploy-68fc6554d6-fftdv 1/1    Running  0         52s
javawebbbluedeploy-68fc6554d6-sxgqg 1/1    Running  0         52s
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get svc
NAME      TYPE        CLUSTER-IP  EXTERNAL-IP  PORT(S)
AGE
javawebbapplivesvc LoadBalancer 10.100.87.253 ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com 80:31359/TCP 60s
kubernetes ClusterIP 10.100.0.1    <none>
```

Load balancers (1/1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

<input checked="" type="checkbox"/>	Name	DNS name	State	VPC ID
<input checked="" type="checkbox"/>	ae0fc82729dfa40bd88...	ae0fc82729dfa40bd88bd26...	–	vpc-04028b965c79a1c2f

Load balancer: ae0fc82729dfa40bd88bd26c491601b6

Details | Listeners | Network mapping | Security | Health checks | Target instances

Details

Load balancer type
Classic

Scheme
Internet-facing

Status
2 of 2 instances in service

Hosted zone
ZQSVJUPU6J1EY

<http://ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com/>

← → ↻ 🏠

⚠ Not secure ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com

📱

🔍 New Tab

🌐 Where should finge...

🌐 http://3.96.216.255/

🗑 Move or copy cells...

🔍 Email Finder hunter.io

🛡 Heroicons

👤 loghmanb/c

Product Management System View Products Add Product Update Product Delete Product


```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat green-deployment.yml
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebgreendeploy
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: java-web-app
      version: v2
      color: green
  template:
    metadata:
      labels:
        app: java-web-app
        version: v2
        color: green
    spec:
      containers:
        - name: javawebappcontainer
          image: hacker123shiva/springbt-in-docker:latest
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
```

...

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat pre-pod.yml
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: javaprepodsvc
spec:
  type: NodePort
  selector:
    app: javawebapp
    color: green
  ports:
    - port: 80
      targetPort: 8080
      protocol: TCP
      nodePort: 31785 # Optional: remove this if you want auto-assign port
```

...

```
ubuntu@ip-172-31-9-165:~/blue-green-model$
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f green-deployment.yml
deployment.apps/javawebgreendeploy created
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f pre-pod.yml
service/javaprepodsvc created
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
jawawebbbluedeploy-68fc6554d6-fftdv	1/1	Running	0	14m

```

javawebbbluedeploy-68fc6554d6-sxgqg 1/1 Running 0 14m
javawebgreendeploy-656f8cf5f4-cn7cb 1/1 Running 0 16s
javawebgreendeploy-656f8cf5f4-jblr5 1/1 Running 0 16s
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)
AGE
javaprepodsvc NodePort      10.100.227.153 <none>
80:31785/TCP  10s
javawebapplivesvc LoadBalancer 10.100.87.253  ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com 80:31359/TCP 14m
kubernetes    ClusterIP     10.100.0.1     <none>
443/TCP
28m

```

Still Blue pods are available for Endusers. Green pods are working in the background

```

ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f green-deployment.yml
deployment.apps/javawebgreendeploy created
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f pre-pod.yml
service/javaprepodsvc created
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get pods
NAME                                READY    STATUS    RESTARTS    AGE
javawebbbluedeploy-68fc6554d6-fftdv 1/1      Running   0            14m
javawebbbluedeploy-68fc6554d6-sxgqg 1/1      Running   0            14m
javawebgreendeploy-656f8cf5f4-cn7cb 1/1      Running   0            16s
javawebgreendeploy-656f8cf5f4-jblr5 1/1      Running   0            16s
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
javaprepodsvc NodePort      10.100.227.153 <none>
80:31785/TCP 10s
javawebapplivesvc LoadBalancer 10.100.87.253  ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com 80:31359/TCP 14m
kubernetes    ClusterIP     10.100.0.1     <none>
443/TCP
28m
ubuntu@ip-172-31-9-165:~/blue-green-model$

```

```

ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f green-deployment.yml
deployment.apps/javawebgreendeploy created
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f pre-pod.yml
service/javaprepodsvc created
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get pods
NAME                                READY    STATUS    RESTARTS    AGE
javawebbbluedeploy-68fc6554d6-fftdv 1/1      Running   0            14m
javawebbbluedeploy-68fc6554d6-sxgqg 1/1      Running   0            14m
javawebgreendeploy-656f8cf5f4-cn7cb 1/1      Running   0            16s
javawebgreendeploy-656f8cf5f4-jblr5 1/1      Running   0            16s
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
javaprepodsvc NodePort      10.100.227.153 <none>
80:31785/TCP 10s
javawebapplivesvc LoadBalancer 10.100.87.253  ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com 80:31359/TCP 14m
kubernetes    ClusterIP     10.100.0.1     <none>
443/TCP
28m
ubuntu@ip-172-31-9-165:~/blue-green-model$

```

How do we access NodePort?

Which service file I got to make changes now so Green deployment goes live instead of Blue?
live-service.yml

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat live-service.yml
```

apiVersion: v1

kind: Service

metadata:

name: javawebapplivesvc

spec:

type: LoadBalancer

selector:

app: java-web-app # Matches the app

color: blue # Sends traffic to the blue pods

color: green

```
ports:
  - port: 80
    targetPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f live-service.yml
service/javawebapplivesvc configured
```

Whatever seconds it takes to re-apply that's the only downtime we have here

Updated Docker image in green deployment so we can see which application goes live

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat green-deployment.yml
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebgreendeploy
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: java-web-app
      version: v2
      color: green
  template:
    metadata:
      labels:
        app: java-web-app
        version: v2
        color: green
    spec:
      containers:
        - name: javawebappcontainer
          image: jmalloc/echo-server
          imagePullPolicy: Always
          ports:
            - containerPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f green-deployment.yml
deployment.apps/javawebgreendeploy configured
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f live-service.yml
service/javawebapplivesvc unchanged
```

Again I go to the same LoadBalancer DNS

<http://ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com/>

← → ↻ 🏠 ⚠ Not secure ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com

🗪 New Tab 🔍 Where should finge... 🌐 http://3.96.216.255/ 🗑 Move or copy cells... 📧 Email Finder hunter.io 🛡 Heroicons 🗨 loghmanb/daily-co... 🖨

Request served by javawebgreendeploy-6bb7bf9f95-gcszv

GET / HTTP/1.1

Host: ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com
 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.7
 Accept-Encoding: gzip, deflate
 Accept-Language: en-US,en;q=0.9
 Cache-Control: max-age=0
 Connection: keep-alive
 Upgrade-Insecure-Requests: 1
 User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/136.0.0.0 Safari/537.36

Going back to Blue

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat live-service.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: javawebapplivesvc
```

```
spec:
```

```
  type: LoadBalancer
```

```
  selector:
```

```
    app: java-web-app # Matches the app
```

```
    color: blue # Sends traffic to the blue pods
```

```
    #color: green
```

```
  ports:
```

```
    - port: 80
```

```
      targetPort: 8080
```

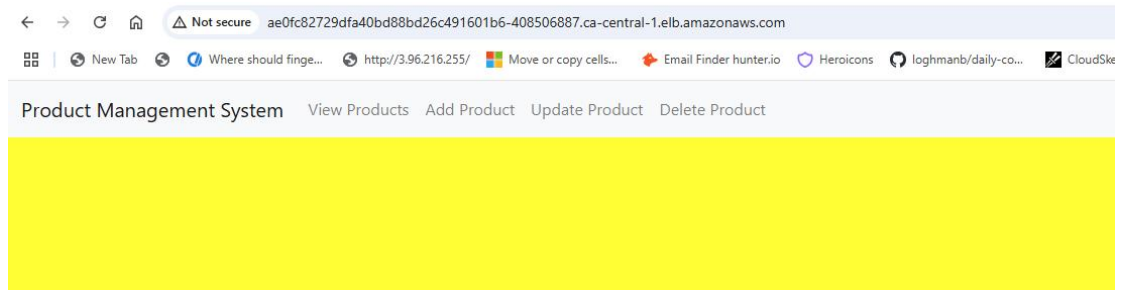
```
...
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat live-service.yml
ubuntu@ip-172-31-9-165:~/blue-green-model$ cat live-service.yml
---
apiVersion: v1
kind: Service
metadata:
  name: javawebapplivesvc
spec:
  type: LoadBalancer
  selector:
    app: java-web-app # Matches the app
    color: blue # Sends traffic to the blue pods
    #color: green
  ports:
    - port: 80
      targetPort: 8080
...
```

```
ubuntu@ip-172-31-9-165:~/blue-green-model$ kubectl apply -f live-service.yml
service/javawebapplivesvc configured
```

Now it switched back to the Old Blue environment

<http://ae0fc82729dfa40bd88bd26c491601b6-408506887.ca-central-1.elb.amazonaws.com/>



Delete cluster

```
eksctl delete cluster --name my-eks-cluster --region ca-central-1
```